

Super Computing on the go: MapReduce for mobile GPUs



Super Computing on the go: MapReduce for mobile GPUs



Background

- Data Driven Applications
- MapReduce
- GPGPUs

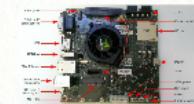
Jetson TK1

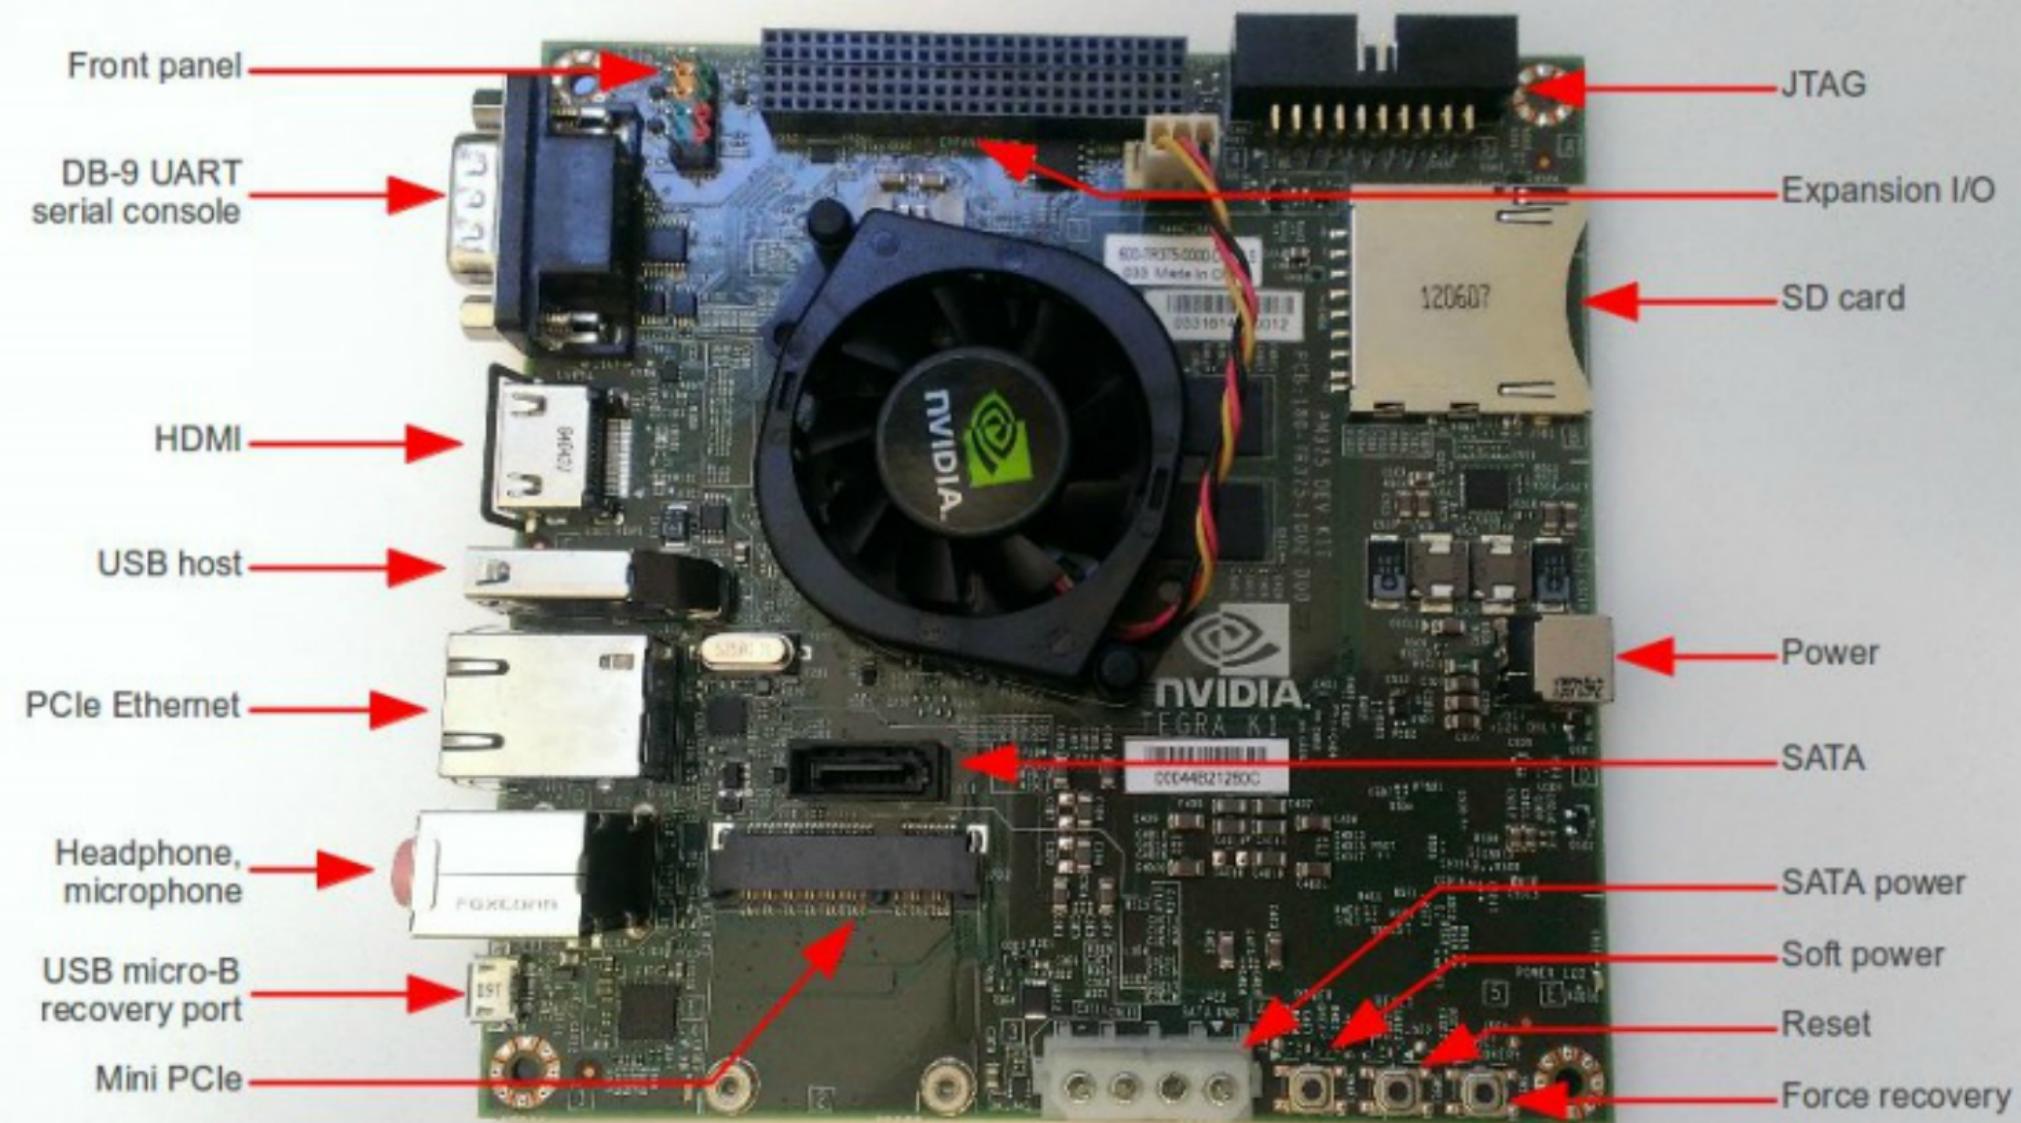
- Tegra K1 - first mobile GPU
- \$192 USD
- 192 cores
- Quad-core ARM processor
- unified memory



Jetson TK1

- Tegra K1 - first mobile GPU
- \$192 USD
- 192 cores
- Quad-core ARM processor
- unified memory





Jetson TK1

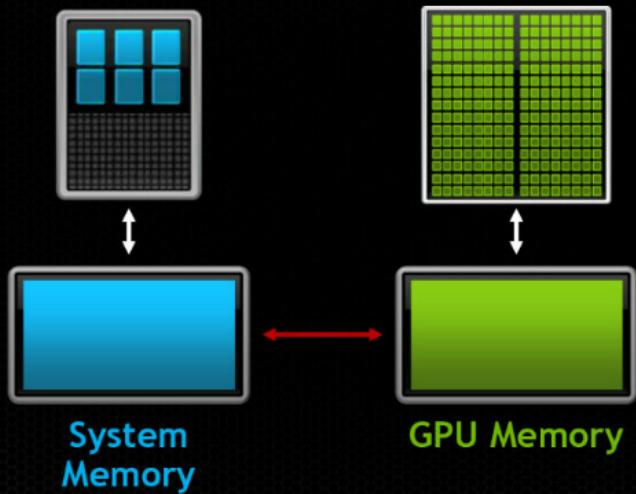
- Tegra K1 - first mobile GPU
- \$192 USD
- 192 cores
- Quad-core ARM processor
- unified memory



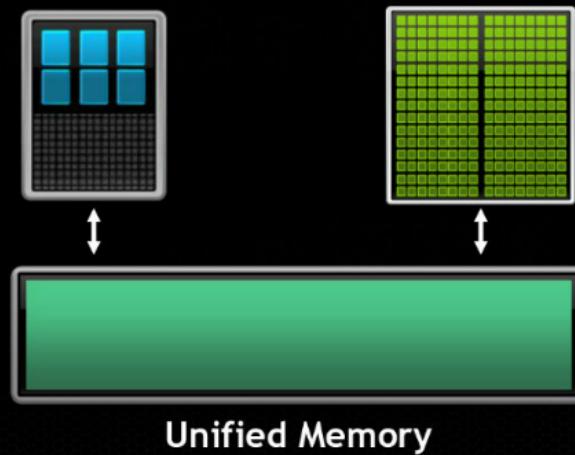
Unified Memory

Dramatically Lower Developer Effort

Developer View Today



Developer View With Unified Memory



Background

MapReduce for Jetson TK1

Features

- Highly parallel computation
- Map Reduce Programming with ease
- Scale-out Architecture
- Mobility
- Super Computing on a power budget
- First of its kind

Proof of Concepts

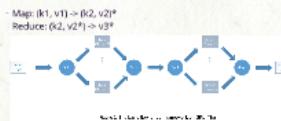
- Hadoop on VM
- Hadoop on Jetson TK1-single node
- MARS on Jetson

Significance

- Mobility
- Power Consumption
- Performance

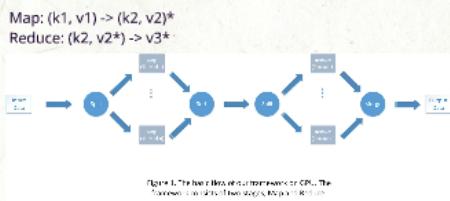
Features

- Highly parallel computation
- Map Reduce Programming with ease
- Scale-out Architecture
- Mobility
- Super Computing on a power budget
- First of its kind



Features

- Highly parallel computation
- Map Reduce Programming with ease
- Scale-out Architecture
- Mobility
- Super Computing on a power budget
- First of its kind



Map: $(k_1, v_1) \rightarrow (k_2, v_2)^*$

Reduce: $(k_2, v_2^*) \rightarrow v_3^*$

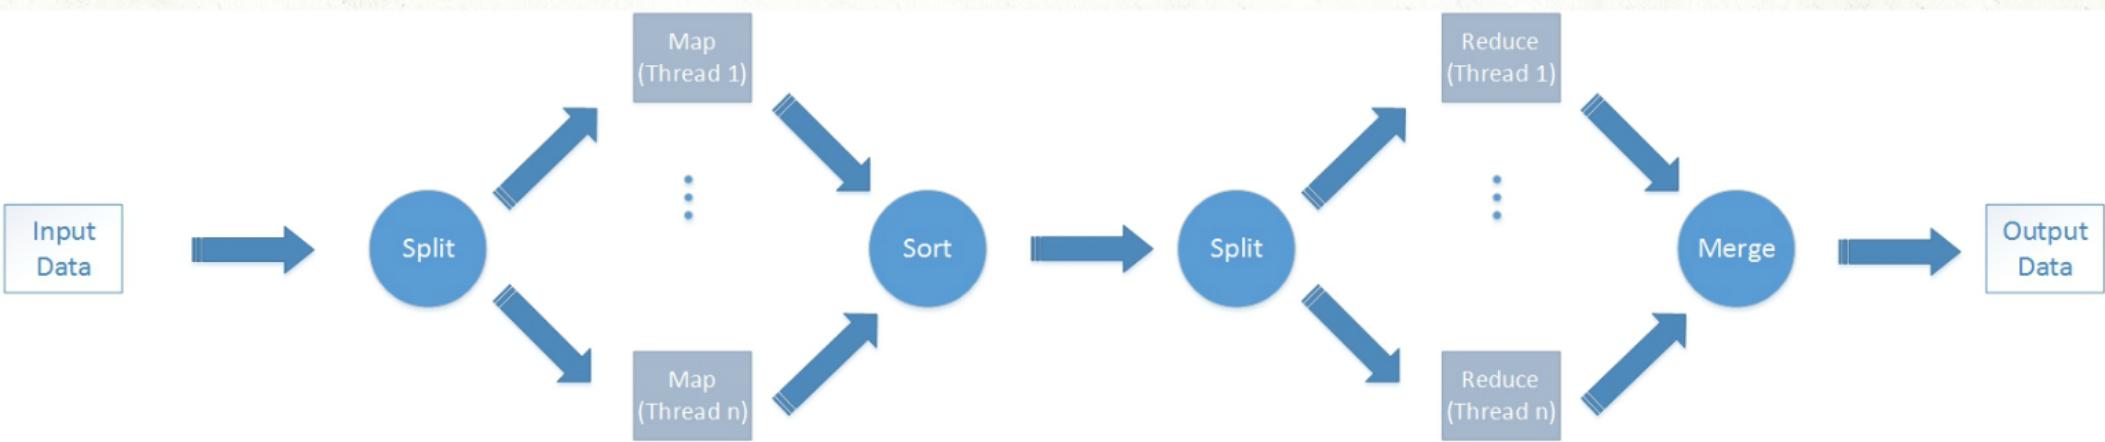


Figure 1. The basic flow of our framework on GPU. The framework consists of two stages, Map and Reduce.

Significance

- Mobility
- Power Consumption
- Performance

Proof of Concepts

- Hadoop on VM
- Hadoop on Jetson TK1-single node
- MARS on Jetson

Literature Review

- Remote Access - modmpi.com/blog
- Big Data Analytics - "World's largest artificial neural networks with GPUs"
- MPI4Py performance improvements
- CUDA-Dark Parallel Template Library
- except a header-only C++11 library for constructing and parsing OpenSoundControl packets
- MPI Cluster System
 - OpenMPI Cluster API Developer Zone
 - CUDA Zone - Nvidia Developers
- History and Evolution of GPU Architecture
- GPU Clusters for High-Performance Computing
- Current needs for using GPUs for accelerating CT reconstruction quality, performance, and timing
- MR-Espresso NVIDIA Tegra K1
- GPGPU Processing in CUDA Architecture
- Multi-GPU Accelerated Image Processing on Large Clusters
- Hadoop and HDFS - Clusters
- Jetson TK1 Embedded Development Kit - NVIDIA
- CUDA 6 Unified Memory explained - Blog - StreamComputing

Literature Review

- Remote Access - modmypi.com/blog
- Big Data Analytics - "World's largest artificial neural networks with GPUs"
- NVIDIA performance primitives
- CUDA Data Parallel Primitives Library
- oscpp a header-only C++11 library for constructing and parsing OpenSoundControl packets
- MPI Cluster System
- OpenCL technology - Intel Developer Zone
- CUDA Zone - NVidia Developer
- History and Evolution of GPU Architecture
- GPU Clusters for High-Performance Computing
- Current and next-generation GPUs for accelerating CT reconstruction: quality, performance, and tuning
- Whitepaper NVIDIA Tegra K1
- GPGPU Processesing in CUDA Architecture
- MapReduce: Simplified Data Processing on Large Clusters
- Hadoop and HDFS: - Cloudera
- Jetson TK1 Embedded Development Kit | NVIDIA
- CUDA 6 Unified Memory explained - Blog - StreamComputing

Literature Review

- Remote Access - modmypi.com/blog
- Big Data Analytics - "World's largest artificial neural networks with GPUs"
- NVIDIA performance primitives
- CUDA Data Parallel Primitives Library
- oscpp a header-only C++11 library for constructing and parsing OpenSoundControl packets
- MPI Cluster System
- OpenCL technology - Intel Developer Zone
- CUDA Zone - NVidia Developer
- History and Evolution of GPU Architecture
- GPU Clusters for High-Performance Computing
- Current and next-generation GPUs for accelerating CT reconstruction: quality, performance, and tuning
- Whitepaper NVIDIA Tegra K1

- MRI Cluster System

- OpenCL technology - Intel Developer Zone
- CUDA Zone - NVidia Developer
- History and Evolution of GPU Architecture
- GPU Clusters for High-Performance Computing
- Current and next-generation GPUs for accelerating CT reconstruction: quality, performance, and tuning
- Whitepaper NVIDIA Tegra K1
- GPGPU Processing in CUDA Architecture
- MapReduce: Simplified Data Processing on Large Clusters
- Hadoop and HDFS: - Cloudera
- Jetson TK1 Embedded Development Kit | NVIDIA
- CUDA 6 Unified Memory explained - Blog - StreamComputing

Challenges

- Released lately
- Modified version of CUDA
- Linux for Tegra
- Memory Architecture
- Concurrent Reads and Writes
- Efficient Load Balancing Scheme
- Low Synchronization Overhead

- Released lately
- Modified version of CUDA
- Linux for Tegra
- Memory Architecture
- Concurrent Reads and Writes
- Efficient Load Balancing Scheme
- Low Synchronization Overhead

What we've done so far

- Designing - single node
- Coding - single node
- Testing - single node



Designing

GPU

- copy the input data to the GPU
- make Map, Sort, and Reduce kernel calls
- copy the output data back to the CPU
- Each data item is assigned to one GPU thread
- we process many items (a "Chunk") with a single kernel call

CPU

- Take a set of indivisible work units (items)
- Map them, generating key-value pairs
- Sort these pairs by key and Reduce all like keyed pairs and gather the final output

Optimizations

Motto
*focus on reducing communication times
at the expense of more computation time*



Optimizations

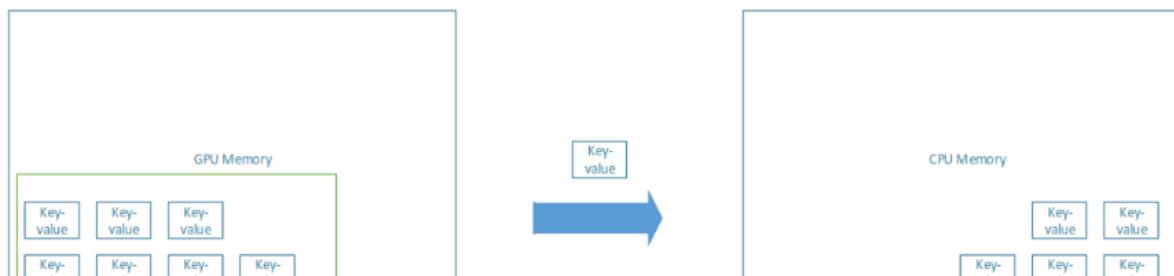
Motto

*focus on reducing communication times
at the expense of more computation time*

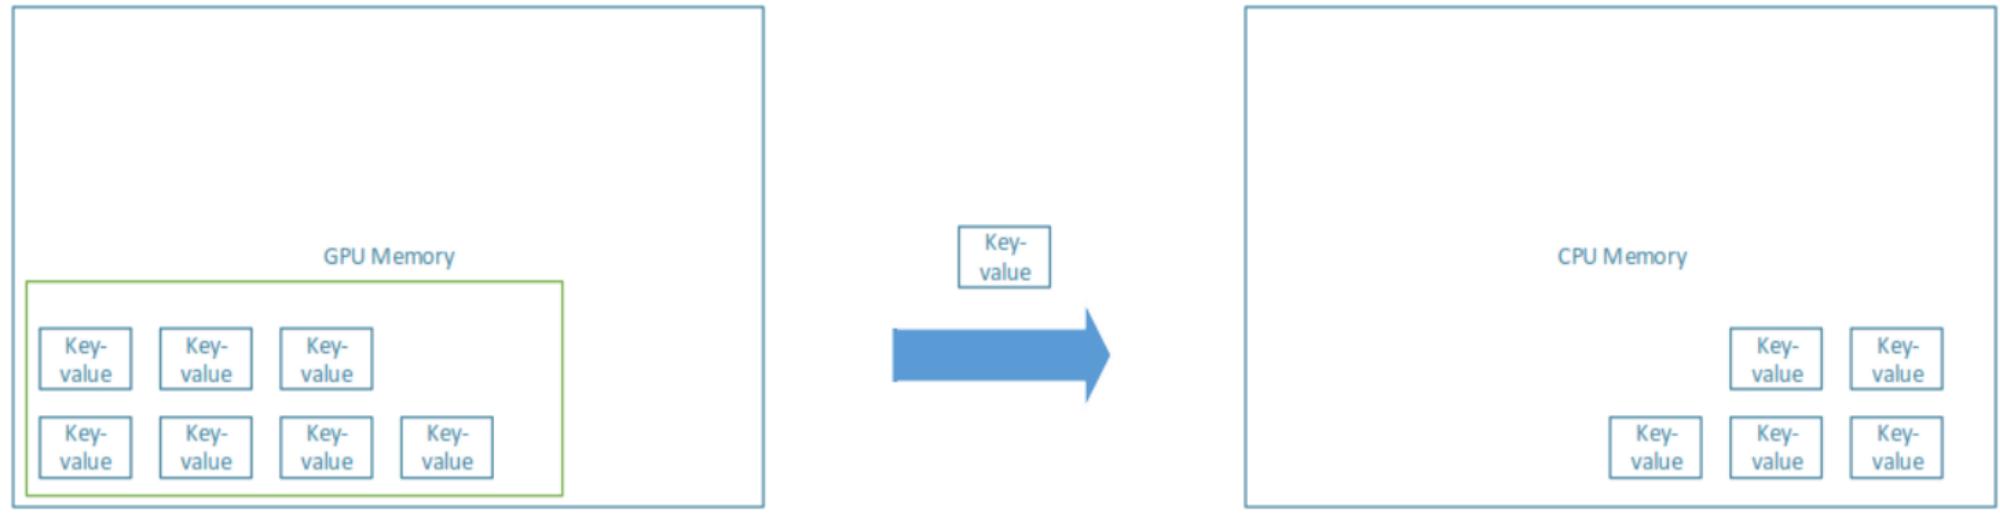
- Chunking
- relaxed mapping, more flexibility and more higher-level operations efficient
- Combine - not new, but needed GPU implementation
- Partial Reduction - final key-value set is large
- Accumulation - final key-value set is small



- Chunking
- relaxed mapping, more flexibility and more higher-level operations efficient
- Combine - not new, but needed GPU implementation
- Partial Reduction - final key-value set is large
- Accumulation - final key-value set is small



Accumulation - final key-value set is s



The polymorphic Map

- Map itself
- Accumulation
- Partial Reduction
- Combination
- one chunk at a time

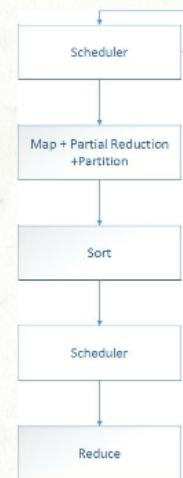
Sort

radix sort from CUDPP (CUDA Data Parallel
Primitives Library)

Reduce

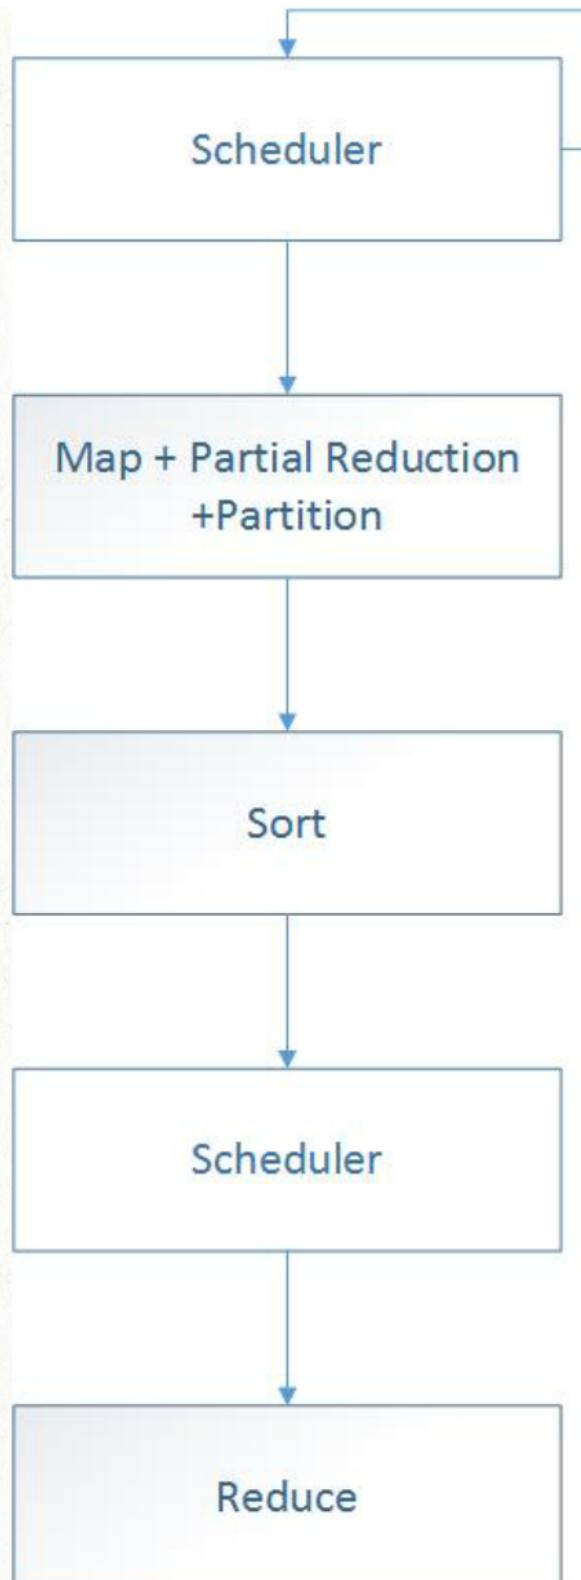
callback to asks how many value sets should be copied for the next reduction

Pipeline Architecture



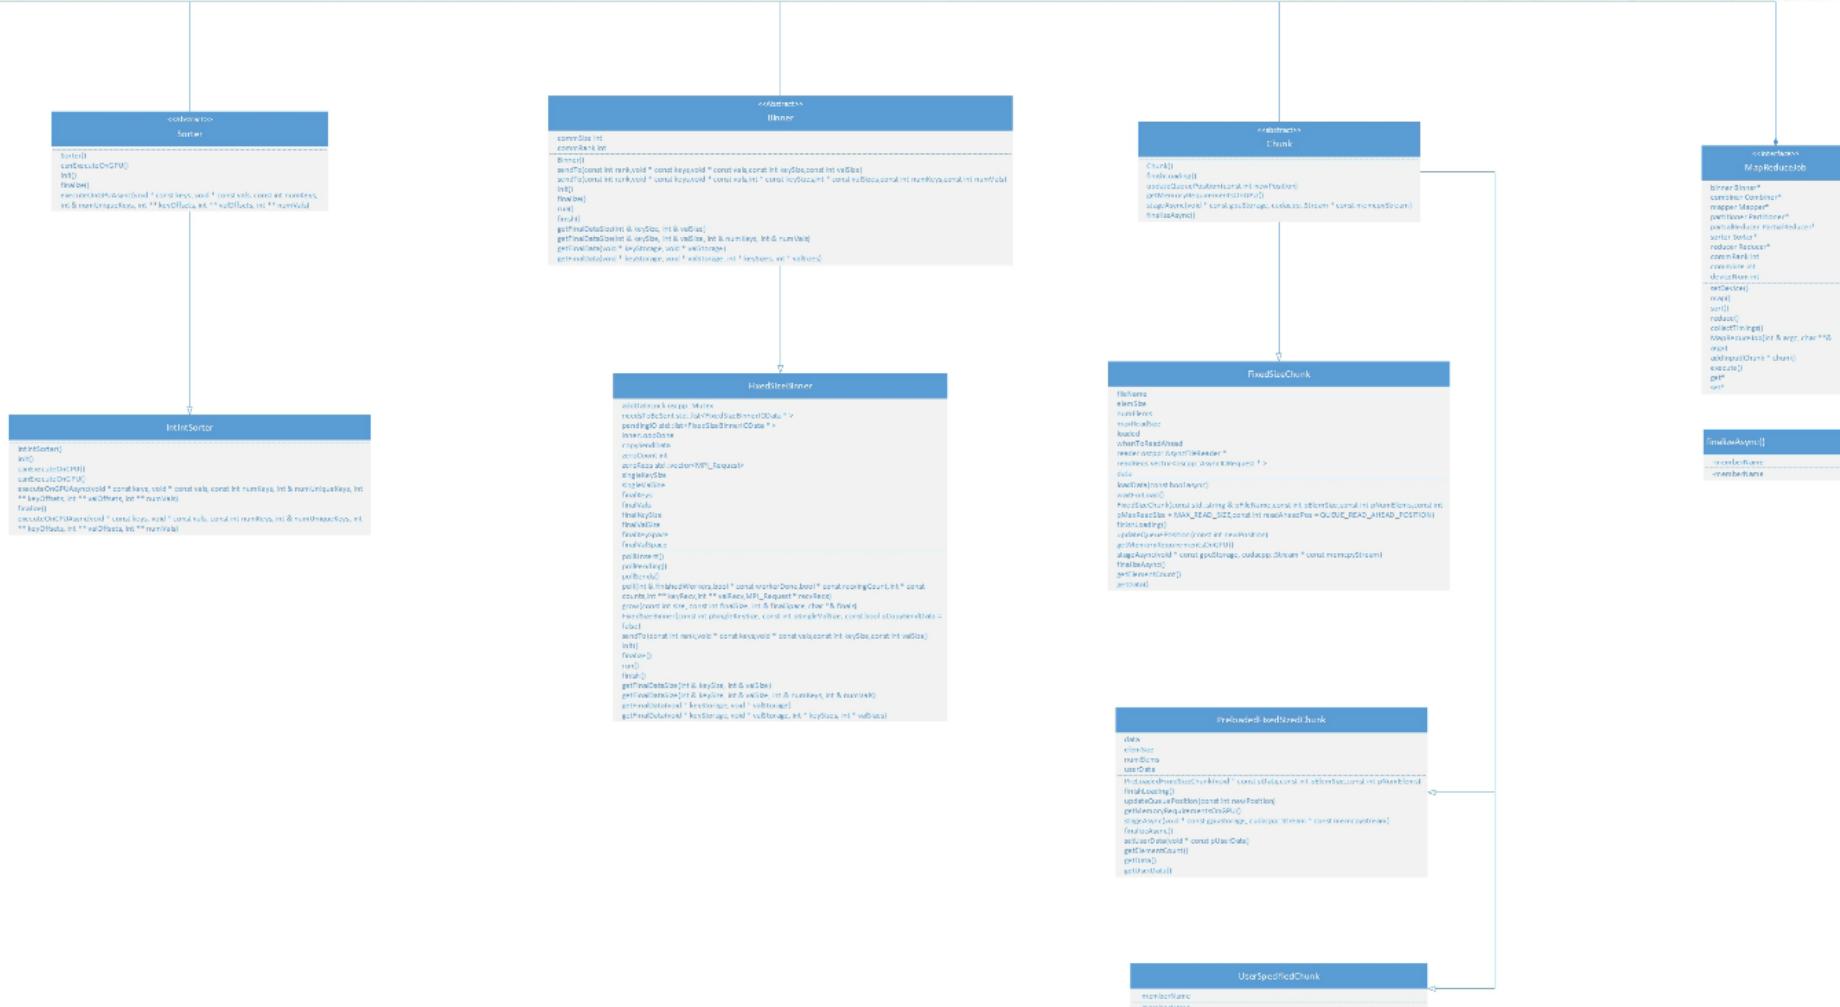
Class diagram

Sequence diagrams

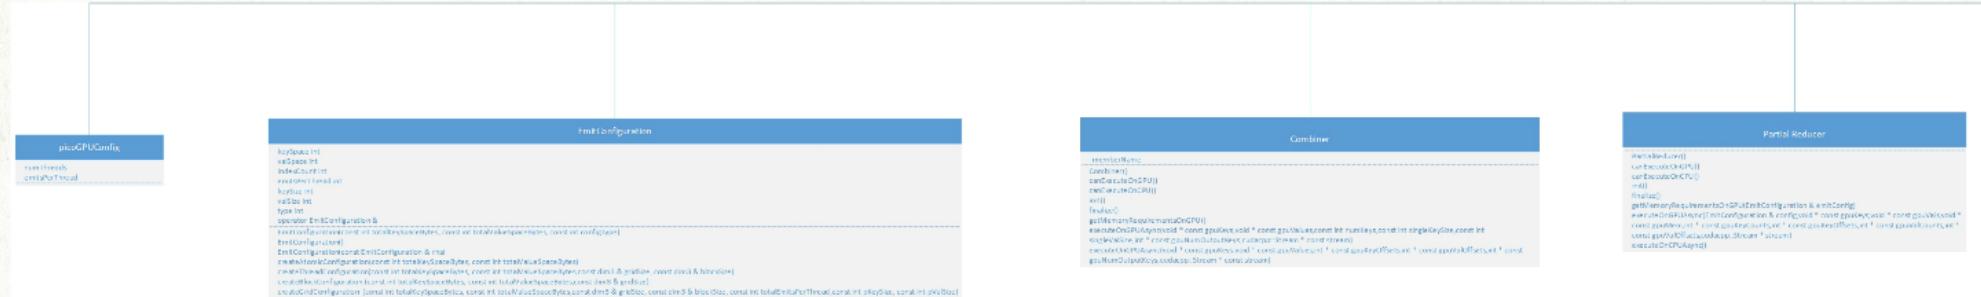


Class diagram



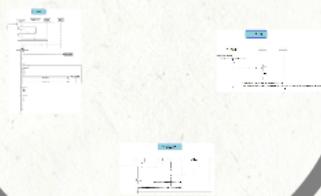




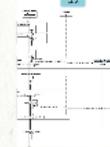


Sequence diagrams

Map



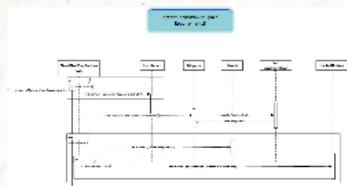
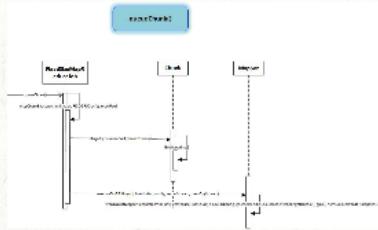
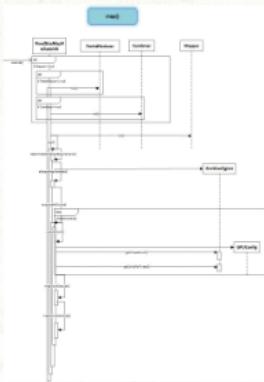
Reduce

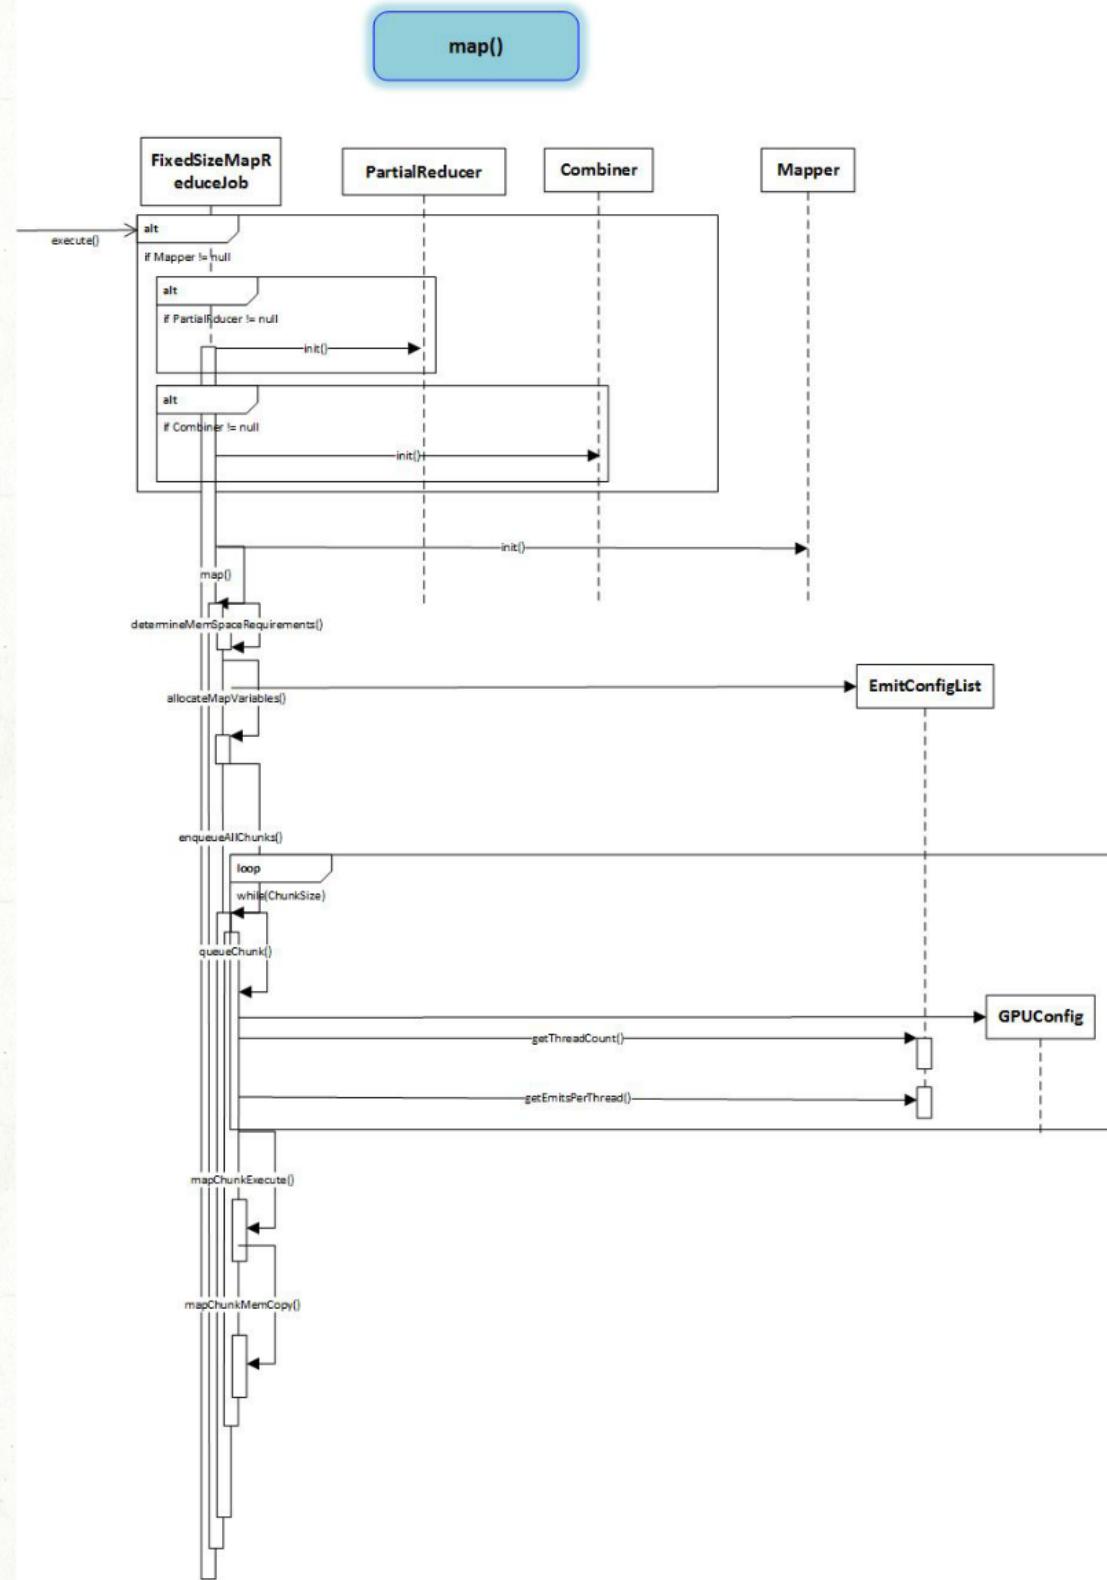


Sort

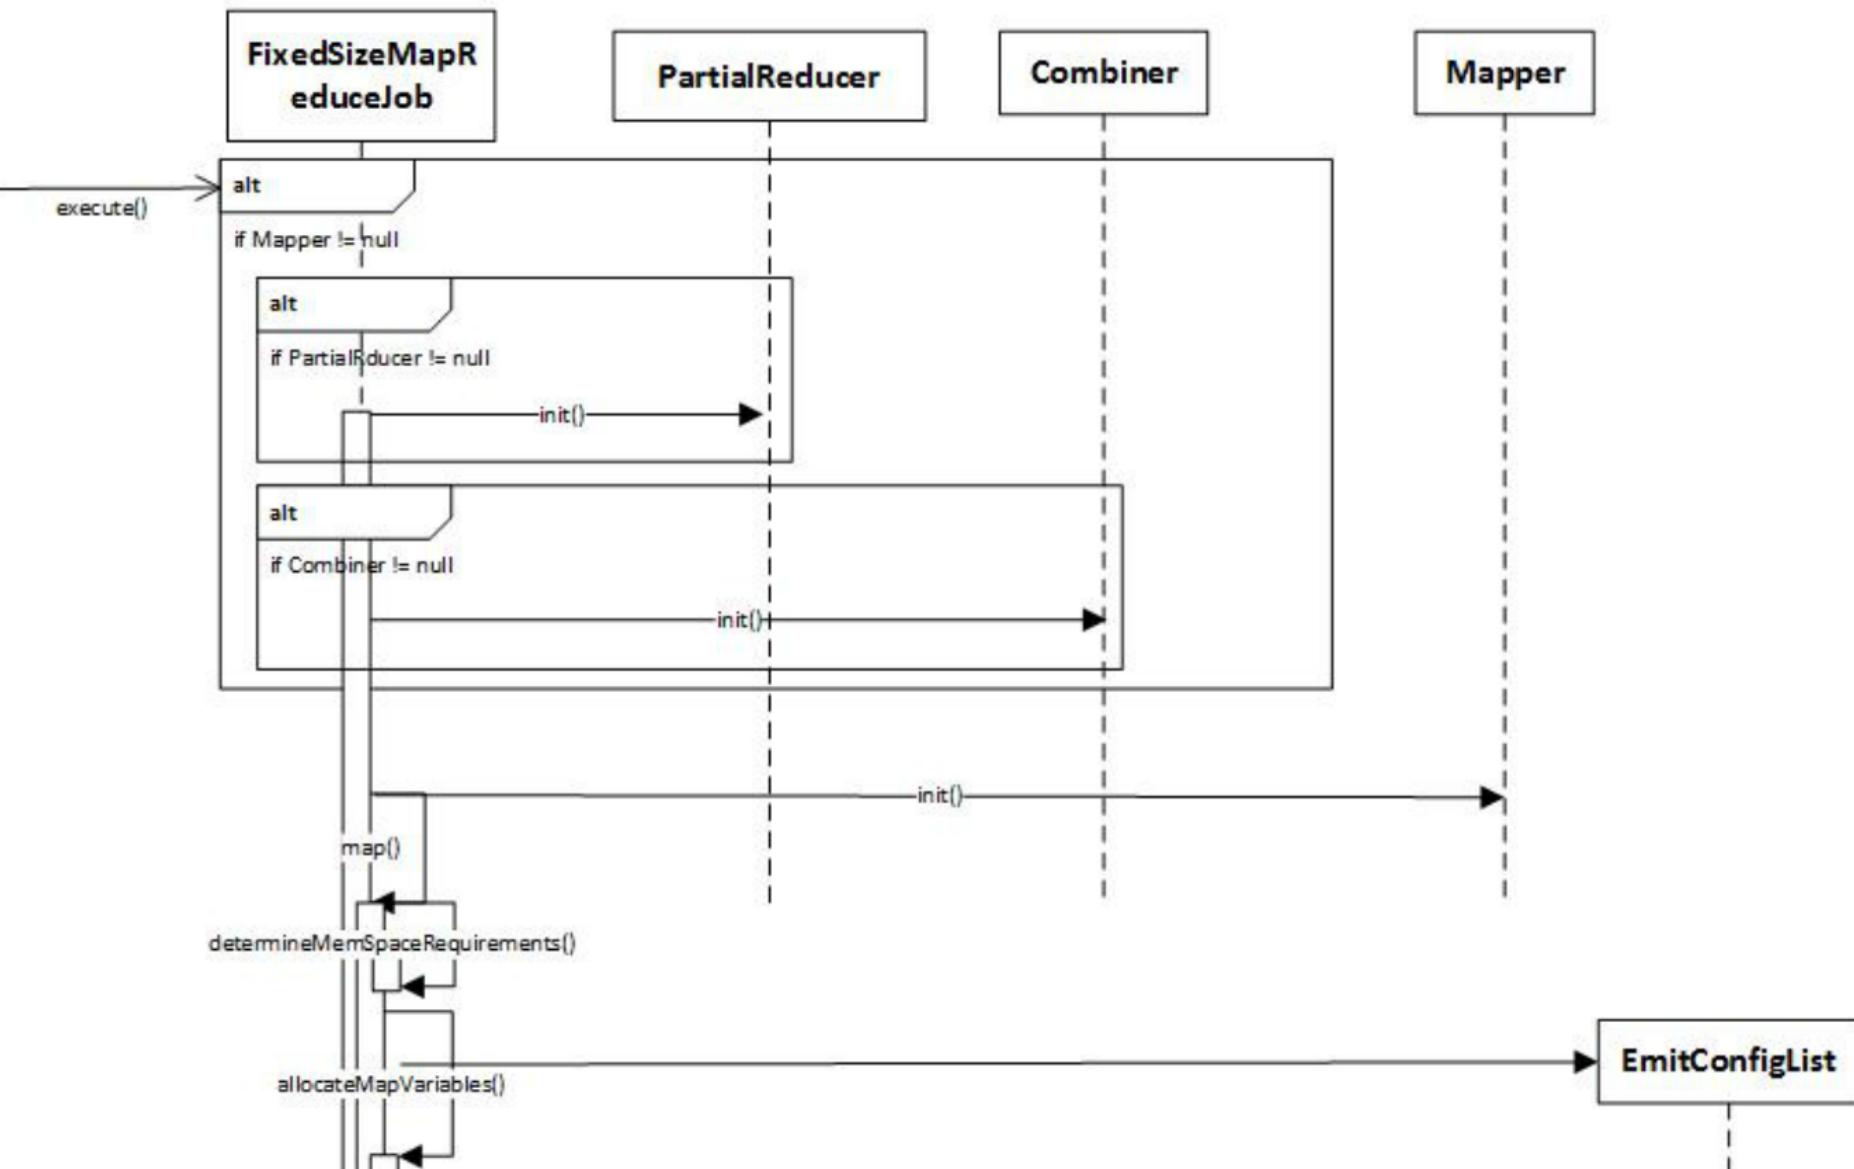


Map



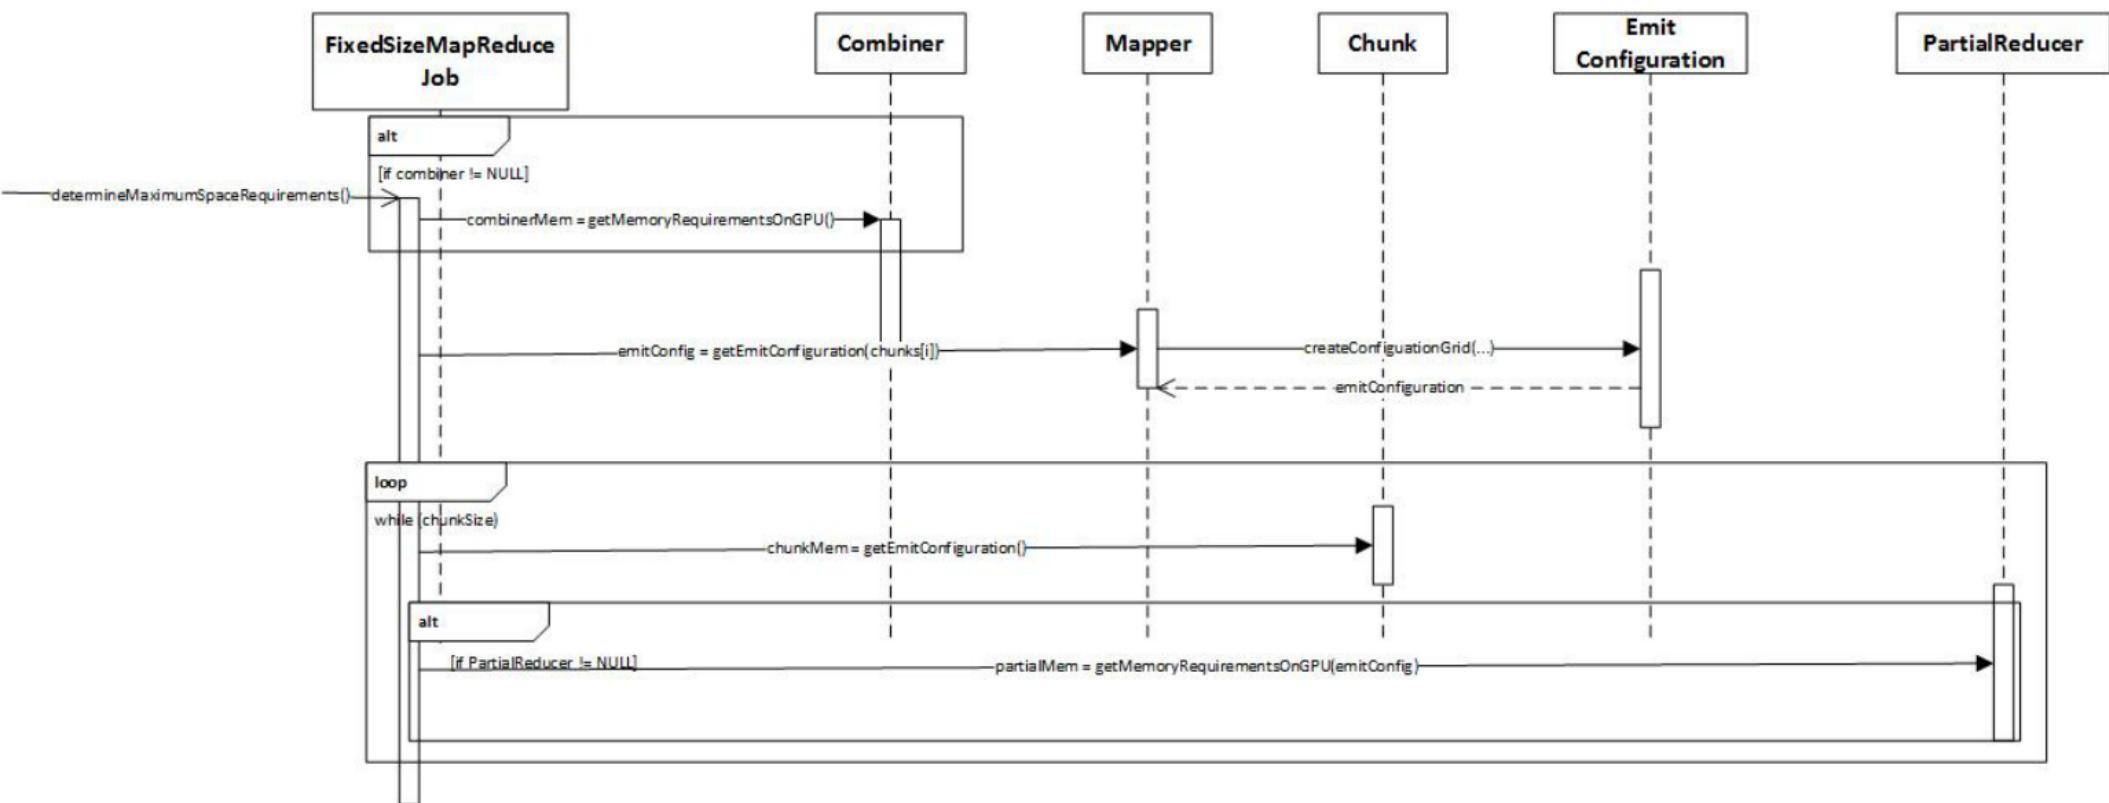


map()

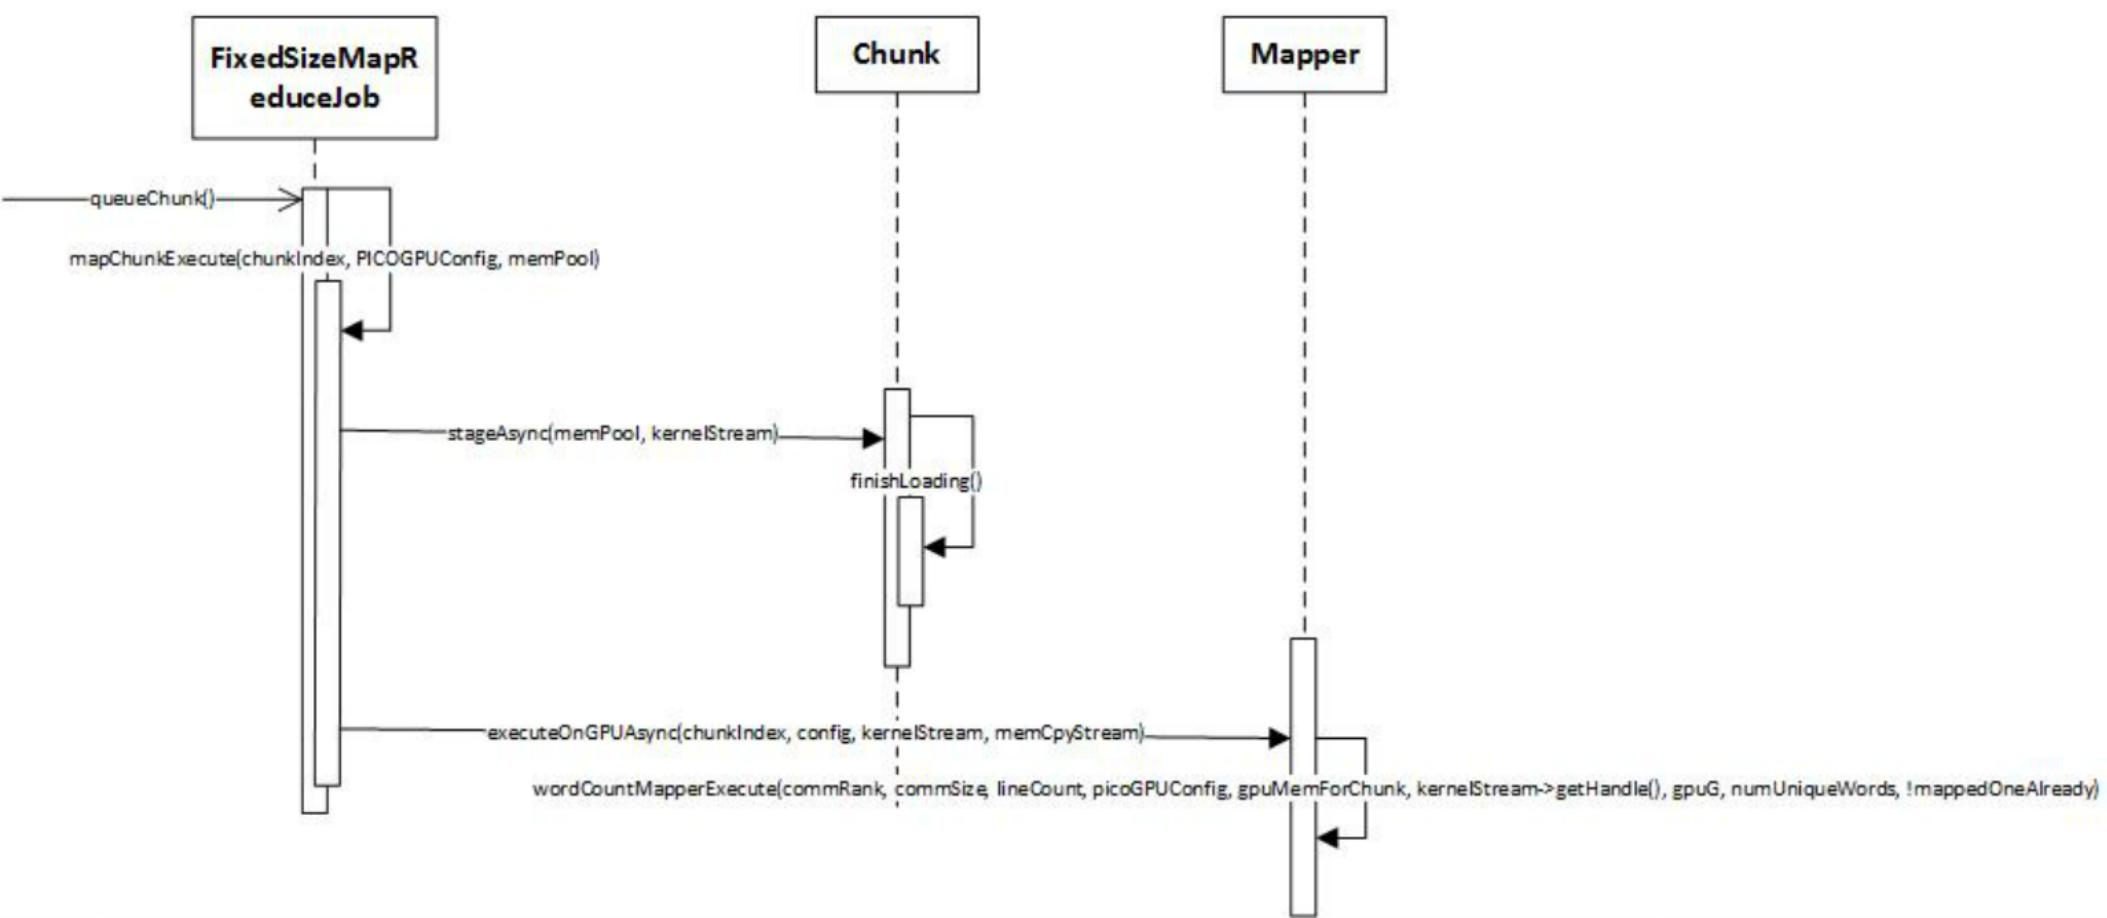




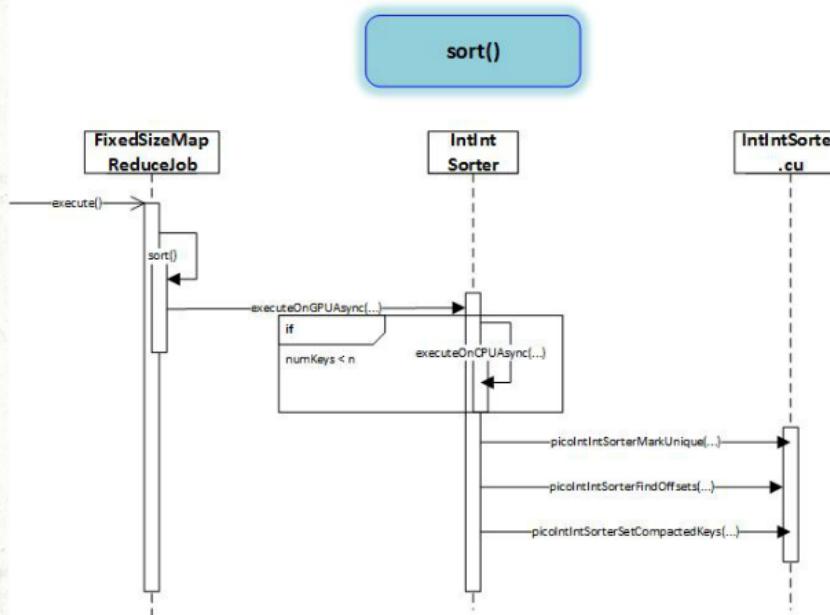
determineMaximumSpace
Requirements()



queueChunk()



Sort

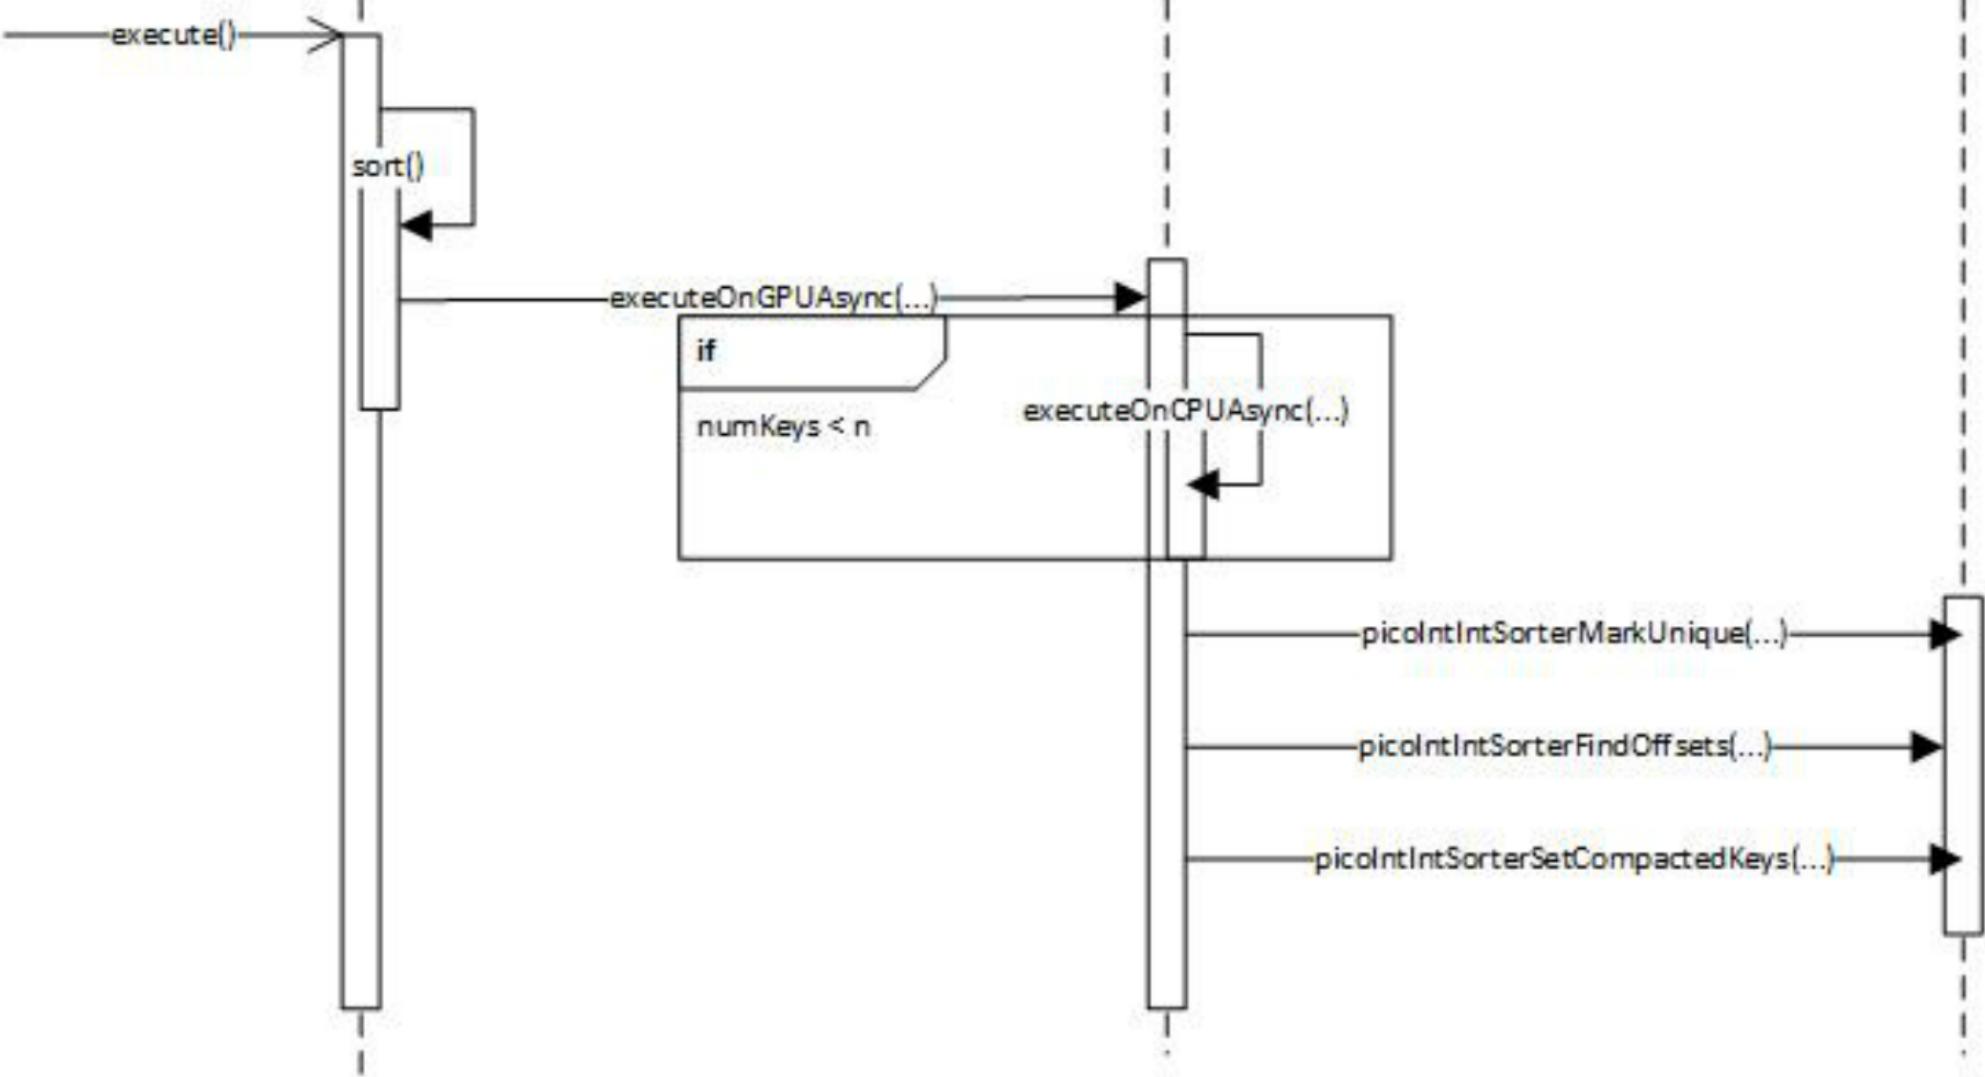


sort()

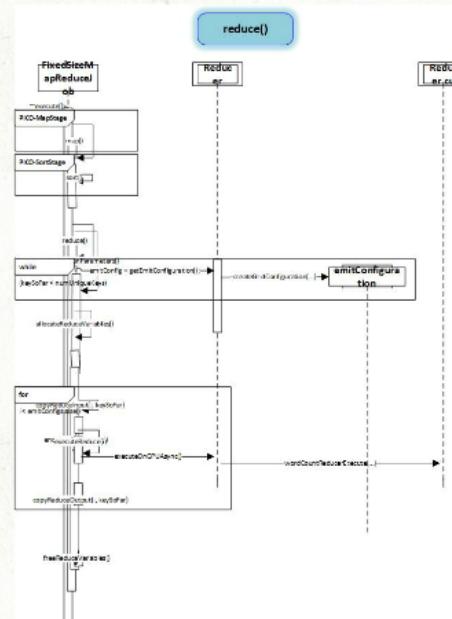
FixedSizeMap
ReduceJob

IntInt
Sorter

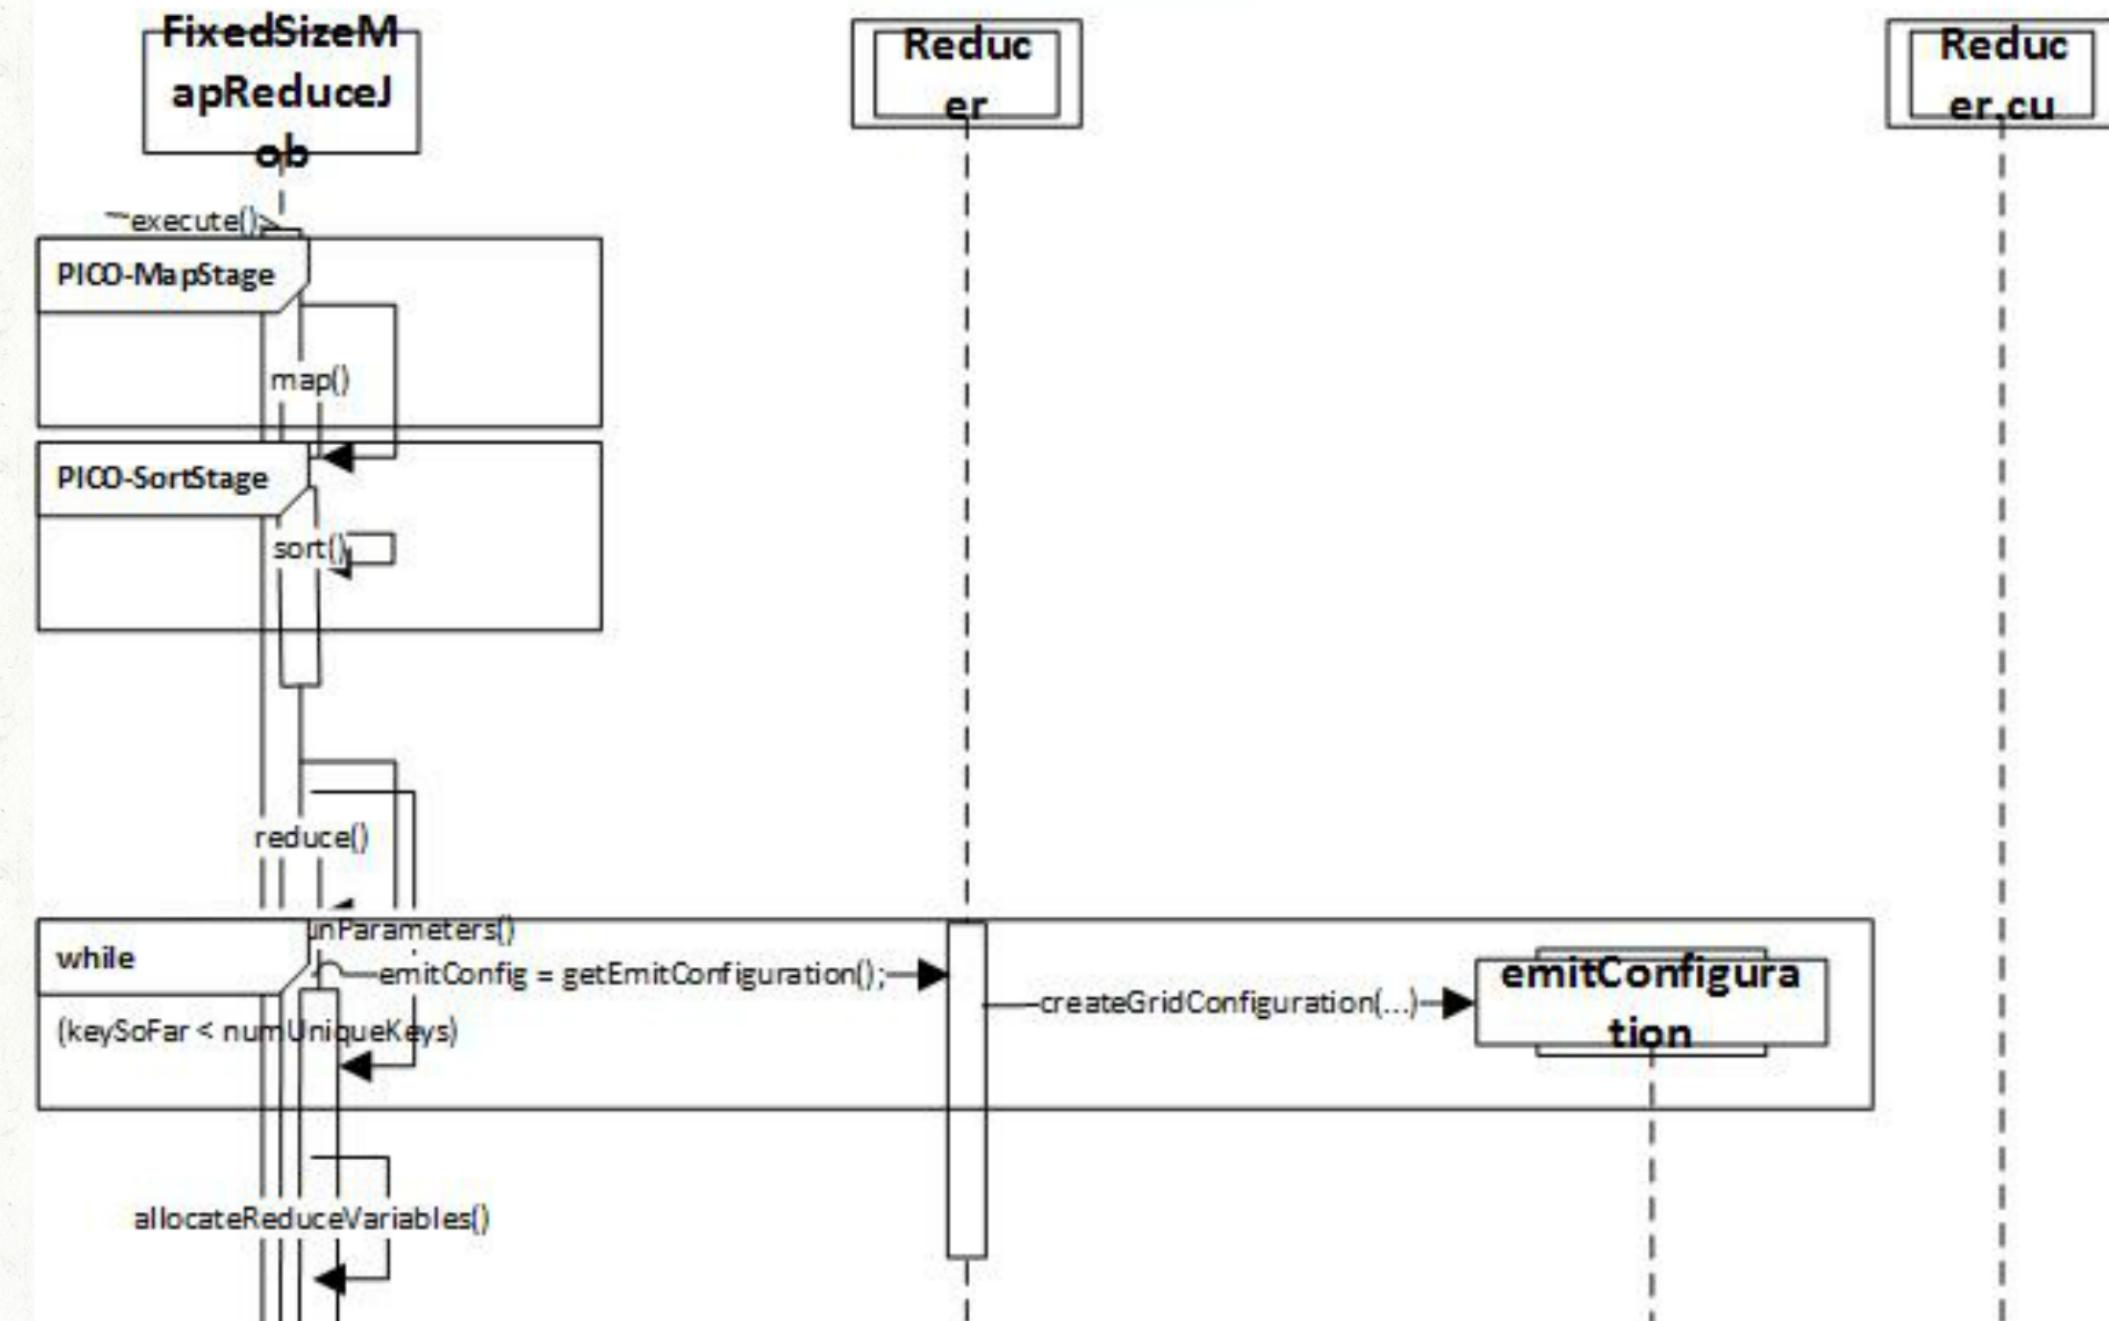
IntIntSorter
.cu



Reduce



reduce()





Results

- Word frequency Benchmark
- 3 subsets of dataset
- 2.47 Speed-up over Mars

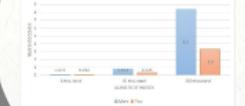
Work Distribution

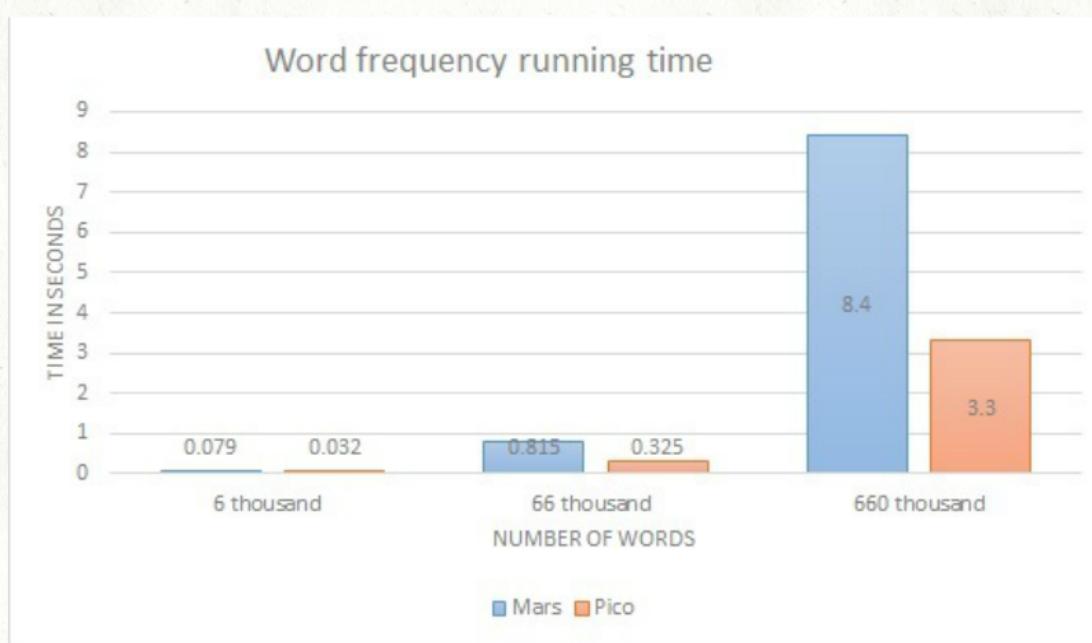


Challenges

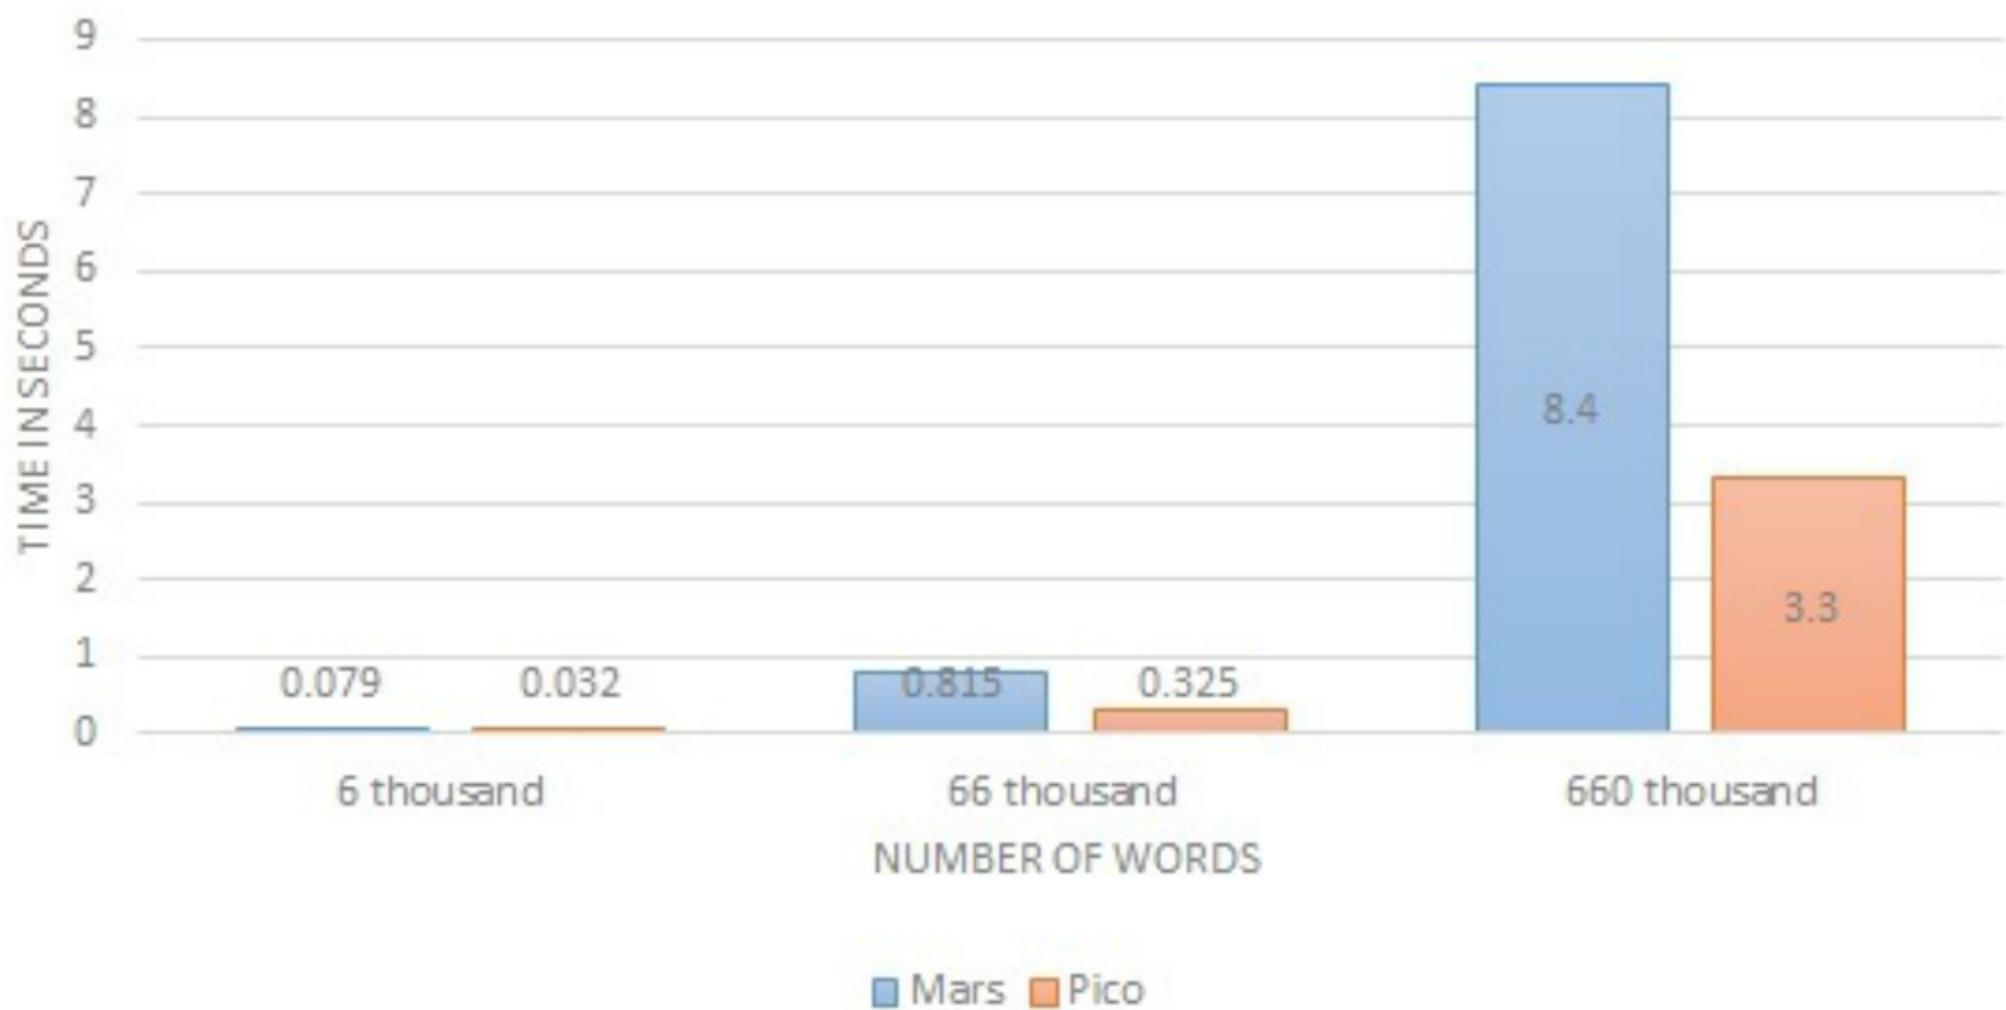
- Released lately
- Modified version of CUDA
- Linux for Tegra
- Memory Architecture
- Concurrent Reads and Writes
- Efficient Load Balancing Scheme
- Low Synchronization Overhead

Word frequency running time





Word frequency running time



Work Distribution

Phase	Team Member
Design	Ahmad Rahman, Abdul Basit, M. Usman Irfan
Implementation-Map	Ahmad Rahman, M. Usman Irfan
Implementation-Sort	Abdul Basit
Implementation-Reduce	Abdul Basit
Testing	Ahmad Rahman, M. Usman Irfan
Documentation	Ahmad Rahman, Abdul Basit

Work Distribution

Phase	Team Member
Design	Ahmad Rahman, Abdul Basit, M. Usman Irfan
Implementation-Map	Ahmad Rahman, M. Usman Irfan
Implementation-Sort	Abdul Basit
Implementation-Reduce	Abdul Basit
Testing	Ahmad Rahman, M. Usman Irfan
Documentation	Ahmad Rahman, Abdul Basit

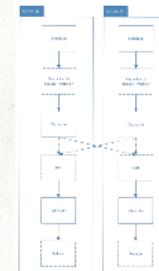
Future Direction

Multilevel Parallelism on
a cluster of Jetson TK1

Hardware

- NVidia Jetson TK1*
- Gigabit Switch
- SATA Hard Drive

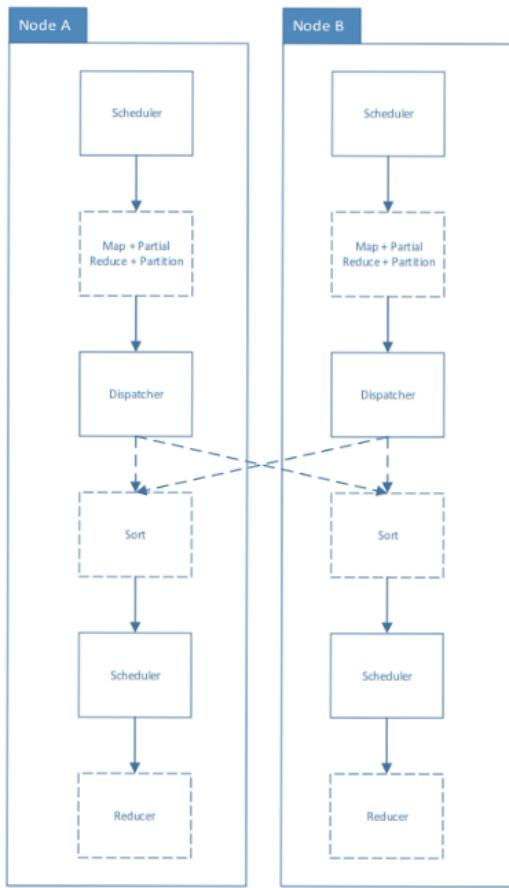
PICO - Multinode

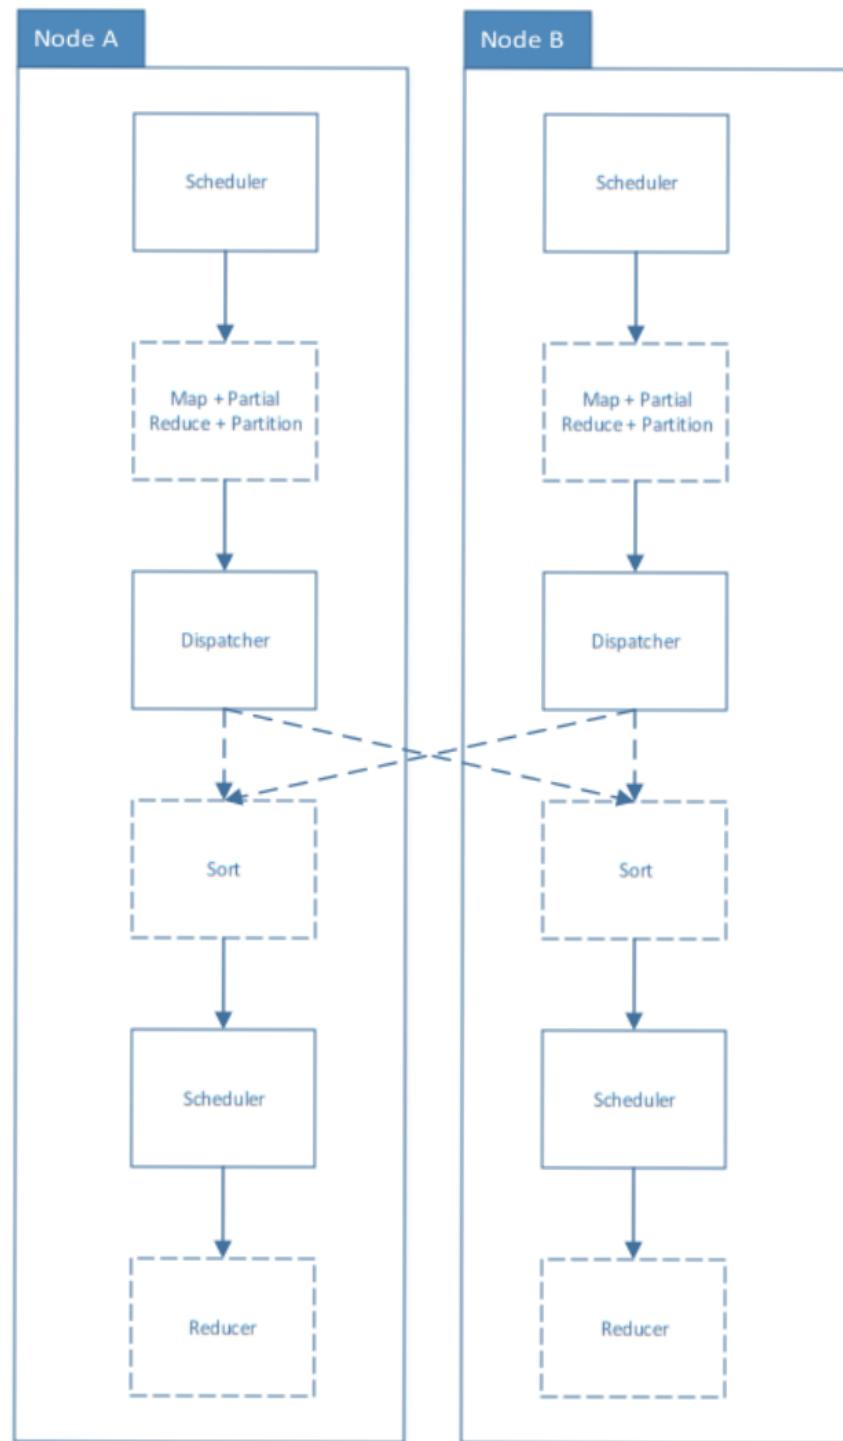


Hardware

- NVidia Jetson TK1*
- Gigabit Switch
- SATA Hard Drive

PICO - Multinode

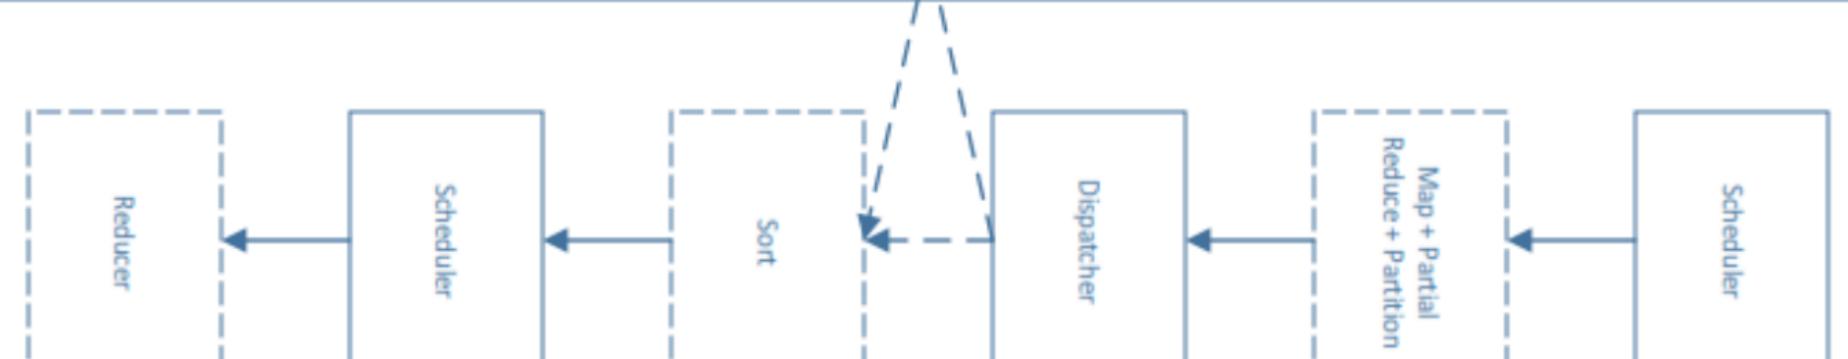




Node A



Node B



References

- <http://blog.kurtosys.com/12-big-facts-about-big-data/#.VBszDBY4S-J>
- <http://www.datanami.com/2014/05/29/hadoop-market-grow-58-2020-report-says/>
- <http://www.developer.com/db/10-facts-about-hadoop.html>
- <http://www.itbusinesseedge.com/blogs/integration/five-facts-for-understanding-big-data-and-hadoop.html>
- elinux.org/Jetson_TK1
- jetsonhacks.com/2015/06/03/battery-power-for-nvidia-jetson-tk1/
- infosys.com - A successful data center migration

Super Computing on the go: MapReduce for mobile GPUs

